

Chapter 1

Computer Abstractions and Technology

What is a Computer ?

- ... ???

What is a Computer ?

Can you solve

$$7x^3 + 5x^2 + 3x + 1 = 0$$

???

What is a Computer ?

Can you solve

$$7x^3 + 5x^2 + 3x + 1 = 0$$

???

If you do not know how to solve a problem,
you cannot tell the computer how to do it !!

What is a Computer ?

- A tool invented to help us do calculations
- Computer is capable of basic logic and arithmetic,
... and nothing else !!!
- Computers are programmed according to algorithms that solve problems ...


What is a Computer ?

- Computers are conceived/designed with an application domain in mind.
- “General Purpose Computers” have wide range of applications.
- The quest for speed leads to special designs: e.g. DSP, VLIW, SuperScalar.

A Fact of Life

- The faster the computer, the more special purpose it becomes ...
- RISC is the right compromise for general purpose computers

A Little History



First (1938-1953): Electronic Numerical Integrator & Computer (ENIAC)
Electronic Discrete Variable Automatic Computer (EDVAC) ←TST Stored Program
Assembly Language, single user, fixed point arithmetic, CPU assisted I/O

Second (1955-1964): TRADIC (Bell Labs), Stretch (IBM 7030, inst. lookahead & error correction),
IBM 7090, CDC 1604, Univac LARC
HLL; FORTRAN (1956), Cobol (1959), Algol (1960),
Compilers, Libraries, Batch, Monitor, Floating Point,
I/O Processors, Multiplexed Memory Access

Third (1965-1974): IBM 360/370, CDC 6600, TI ASC, PDP-8,
ILLIAC 4 (1968) ←48x48 Mesh Connected Parallel Computer)
Microprogramming, Cache, Multiprogramming, Times-shared OS,
Intelligent Compilers, Virtual Memory

Fourth (1975-1990): VAX 9000, Cray X-MP, IBM 3090, BBN TC2000
MPP (Goodyear/NASA)
HLL for Scalar & Vector Data, Vectorizing Compilers,
Languages and Environments for Parallel Processing

Adapted from K. Hwang, *Advanced Computer Architecture: Parallelism, Scalability, Programmability*, McGraw-Hill, Inc., Reading, 1993.


Technology Drives Design

- Computers designed using the latest technological advances.
- Designs always tend to provide best balance of technologies for different parts of the computer (CPU, Memory, I/O)

Classes of Computers

- Desktop computers
 - General purpose, variety of software
 - Subject to cost/performance tradeoff
- Server computers
 - Network based
 - High capacity, performance, reliability
 - Range from small servers to building sized
- Embedded computers
 - Hidden as components of systems
 - Stringent power/performance/cost constraints

The Processor Market



The Computer Revolution

- Progress in computer technology
 - Underpinned by Moore's Law
- Makes novel applications feasible
 - Computers in automobiles
 - Cell phones
 - Human genome project
 - World Wide Web
 - Search Engines
- Computers are pervasive

Modern Computer Revolution

- Just like John Von Neumann, David Patterson and John Hennessy revolutionized the computer when they introduced

RISC

Modern Computer Revolution

- Computer Architecture is an art ...
 - No one design method leading to best computer...
- **RISC** is a design method/philosophy ... derived from observation (experimental computer science): *computers execute 20% of their instructions 80% of the time.*

Modern Computer Revolution

§ 1.1
Introduction

- In this course, we shall design a processor based on the RISC design method.
- Processor is MIPS, an actual commercially available processor ...
 - Focus is on design and not on programming. Assumption: you know assembly language programming ...

Brief Outline of course

- Define Performance to have a cost function to evaluate designs
 - Define the general Instruction Set Architecture, for performance
 - Study number crunching and define the arithmetic of the processor
 - Implement the processor:
 - Single Cycle
 - Multi-Cycle
 - Pipelining
- The Memory Hierarchy: Technology rules !!
Advanced topics (time permitting)

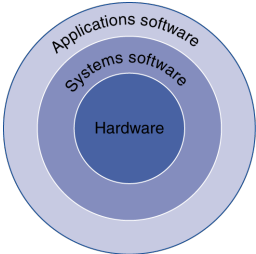
What You Will Learn

- Programs are translated into the machine language (ISA) – [Compilers course]
 - How does the hardware execute them ?
- The hardware/software interface
- What determines program performance
 - And how it can be improved
- How hardware designers improve performance: CPU design – Memory hierarchy design
- Balancing all the components of a system: CPU-Memory-I/O
- What is parallel processing

Understanding Performance

- Algorithm
 - Determines number of operations executed
- Programming language, compiler, architecture
 - Determine number of machine instructions executed per operation
- Processor and memory system
 - Determine how fast instructions are executed
- I/O system (including OS)
 - Determines how fast I/O operations are executed

Below Your Program



- Application software
 - Written in high-level language
- System software
 - Compiler: translates HLL code to machine code
 - Operating System: service code
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- Hardware
 - Processor, memory, I/O controllers

§ 1.2 Below Your Program

Levels of Program Code

- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- Assembly language
 - Textual representation of machine instructions
- Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data

High-level language program (in C)

Assembly language program (for MIPS)

Binary machine language program (for MIPS)

```

swap(int v[], int k)
{
    int temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
                    
```

↓
Compiler

```

swap:
    muli $2, $5, 4
    add  $2, $4, $2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31
                    
```

↓
Assembler

```

000000001010000100000000000011000
00000000000110000001100000100001
10001100011000100000000000000000
10001100111100100000000000000100
10101100111100100000000000000000
10101100011000100000000000000100
0000001111100000000000000001000
                    
```

Abstractions


The BIG Picture

- Abstraction helps us deal with complexity
 - Hide lower-level detail
- Instruction set architecture (ISA)
 - The hardware/software interface
- Application binary interface
 - The ISA plus system software interface
- Implementation
 - The details underlying and interface

Defining Performance

§ 1.4 Performance

- Which airplane has the best performance?



Response Time and Throughput

- Response time
 - How long it takes to do a task
- Throughput
 - Total work done per unit time
 - e.g., tasks/transactions/... per hour
- How are response time and throughput affected by
 - Replacing the processor with a faster version?
 - Adding more processors?
- We'll focus on response time for now...

Relative Performance

- Define Performance = $1/\text{Execution Time}$
- "X is n time faster than Y"


$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution time}_Y}{\text{Execution time}_X} = n$$
- Example: time taken to run a program
 - 10s on A, 15s on B
 - $\text{Execution Time}_B / \text{Execution Time}_A = 15\text{s} / 10\text{s} = 1.5$
 - So A is 1.5 times faster than B

Measuring Execution Time

- Elapsed time
 - Total response time, including all aspects
 - Processing, I/O, OS overhead, idle time
 - Determines system performance
- CPU time
 - Time spent processing a given job
 - Discounts I/O time, other jobs' shares
 - Comprises user CPU time and system CPU time
 - Different programs are affected differently by CPU and system performance

CPU Clocking

- Operation of digital hardware governed by a constant-rate clock



- Clock period: duration of a clock cycle
 - e.g., $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$
- Clock frequency (rate): cycles per second
 - e.g., $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$

CPU Time

$$\begin{aligned}\text{CPU Time} &= \text{CPU Clock Cycles} \times \text{Clock Cycle Time} \\ &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}\end{aligned}$$

- Performance improved by
 - Reducing number of clock cycles
 - Increasing clock rate
 - Hardware designer must often trade off clock rate against cycle count

CPU Time Example

- Computer A: 2GHz clock, 10s CPU time
- Designing Computer B
 - Aim for 6s CPU time
 - Can do faster clock, but causes $1.2 \times$ clock cycles
- How fast must Computer B clock be?

$$\text{Clock Rate}_B = \frac{\text{Clock Cycles}_B}{\text{CPU Time}_B} = \frac{1.2 \times \text{Clock Cycles}_A}{6s}$$

$$\begin{aligned}\text{Clock Cycles}_A &= \text{CPU Time}_A \times \text{Clock Rate}_A \\ &= 10s \times 2\text{GHz} = 20 \times 10^9\end{aligned}$$

$$\text{Clock Rate}_B = \frac{1.2 \times 20 \times 10^9}{6s} = \frac{24 \times 10^9}{6s} = 4\text{GHz}$$

Instruction Count and CPI

$\text{Clock Cycles} = \text{Instruction Count} \times \text{Cycles per Instruction}$

$\text{CPU Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$

- Instruction Count for a program
 - Determined by program, ISA and compiler
- Average cycles per instruction
 - Determined by CPU hardware
 - If different instructions have different CPI
 - Average CPI affected by instruction mix

CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA
- Which is faster, and by how much?

$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= 1 \times 2.0 \times 250\text{ps} = 1 \times 500\text{ps}\end{aligned}$$

A is faster...

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= 1 \times 1.2 \times 500\text{ps} = 1 \times 600\text{ps}\end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{1 \times 600\text{ps}}{1 \times 500\text{ps}} = 1.2$$

...by this much

CPI in More Detail

- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- Weighted average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \underbrace{\frac{\text{Instruction Count}_i}{\text{Instruction Count}}}_{\text{Relative frequency}} \right)$$

CPI Example

- Alternative compiled code sequences using instructions in classes A, B, C

Class	A	B	C
CPI for class	1	2	3
IC in sequence 1	2	1	2
IC in sequence 2	4	1	1

- Sequence 1: IC = 5
 - Clock Cycles
= $2 \times 1 + 1 \times 2 + 2 \times 3$
= 10
 - Avg. CPI = $10/5 = 2.0$
- Sequence 2: IC = 6
 - Clock Cycles
= $4 \times 1 + 1 \times 2 + 1 \times 3$
= 9
 - Avg. CPI = $9/6 = 1.5$

Performance Summary

The BIG Picture

$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
 - Algorithm: affects IC, possibly CPI
 - Programming language: affects IC, CPI
 - Compiler: affects IC, CPI
 - Instruction set architecture: affects IC, CPI, T_c